

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
from sklearn.metrics import precision_recall_fscore_support, average_precision_score
from sklearn.metrics import precision_score, recall_score, f1_score, precision_recall_curve
```

```
In [2]: car_data = pd.read_csv("car.csv")
```

```
In [3]: car_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1151 entries, 0 to 1150
Data columns (total 7 columns):
buying      1151 non-null object
maint      1151 non-null object
doors      1151 non-null int64
persons    1151 non-null int64
lug_boot   1151 non-null object
safety     1151 non-null object
class      1151 non-null object
dtypes: int64(2), object(5)
memory usage: 63.0+ KB
```

```
In [4]: car_data.head()
```

Out[4]:

	buying	maint	doors	persons	lug_boot	safety	class
0	vhigh	vhigh	2	2	small	med	unacc
1	vhigh	vhigh	2	2	small	high	unacc
2	vhigh	vhigh	2	2	med	low	unacc
3	vhigh	vhigh	2	2	med	med	unacc
4	vhigh	vhigh	2	2	med	high	unacc

```
In [5]: car_data.describe()
```

```
Out[5]:
```

	doors	persons
count	1151.000000	1151.000000
mean	3.501303	3.000869
std	1.118130	1.000434
min	2.000000	2.000000
25%	3.000000	2.000000
50%	4.000000	4.000000
75%	4.500000	4.000000
max	5.000000	4.000000

```
In [6]: y=car_data['class']
y=y.to_frame()
y.head()
```

```
Out[6]:
```

	class
0	unacc
1	unacc
2	unacc
3	unacc
4	unacc

```
In [7]: X=car_data
X= X[['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety']]
X.head()
```

```
Out[7]:
```

	buying	maint	doors	persons	lug_boot	safety
0	vhigh	vhigh	2	2	small	med
1	vhigh	vhigh	2	2	small	high
2	vhigh	vhigh	2	2	med	low
3	vhigh	vhigh	2	2	med	med
4	vhigh	vhigh	2	2	med	high

```
In [8]: #Applying Train,Test Split
x_train,x_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=32)
```

```
In [9]: combine=[y_train,y_test]
classmapping={'unacc':0,'acc':1,'good':2,'vgood':3}
for dt in combine:
    dt['class']=car_data['class'].map(classmapping)
y_train.head()
```

Out[9]:

	class
921	0
721	0
501	1
432	0
711	0

```
In [10]: combine=[X_train,X_test]
classmapping={'high':0,'low':1,'med':2,'vhigh':3}
for dt in combine:
    dt['buying']=car_data['buying'].map(classmapping)
X_train.head()
```

Out[10]:

	buying	maint	doors	persons	lug_boot	safety
921	1	vhigh	5	2	med	med
721	2	med	2	2	small	high
501	0	med	5	4	big	med
432	0	med	2	2	small	med
711	2	high	5	4	small	med

```
In [11]: combine=[X_train,X_test]
classmapping={'high':0,'low':1,'med':2,'vhigh':3}
for dt in combine:
    dt['maint']=car_data['maint'].map(classmapping)
X_train.head()
```

Out[11]:

	buying	maint	doors	persons	lug_boot	safety
921	1	3	5	2	med	med
721	2	2	2	2	small	high
501	0	2	5	4	big	med
432	0	2	2	2	small	med
711	2	0	5	4	small	med

```
In [12]: combine=[X_train,X_test]
classmapping={'big':0,'small':1,'med':2}
for dt in combine:
    dt['lug_boot']=car_data['lug_boot'].map(classmapping)
X_train.head()
```

```
Out[12]:
```

	buying	maint	doors	persons	lug_boot	safety
921	1	3	5	2	2	med
721	2	2	2	2	1	high
501	0	2	5	4	0	med
432	0	2	2	2	1	med
711	2	0	5	4	1	med

```
In [13]: combine=[X_train,X_test]
classmapping={'high':0,'low':1,'med':2}
for dt in combine:
    dt['safety']=car_data['safety'].map(classmapping)
X_train.head()
```

```
Out[13]:
```

	buying	maint	doors	persons	lug_boot	safety
921	1	3	5	2	2	2
721	2	2	2	2	1	0
501	0	2	5	4	0	2
432	0	2	2	2	1	2
711	2	0	5	4	1	2

```
In [14]: #AdaBoost
clf = AdaBoostClassifier(DecisionTreeClassifier(max_depth=1),
                        algorithm="SAMME",
                        n_estimators=200)
clf.fit(X_train, y_train)
```

```
Out[14]: AdaBoostClassifier(algorithm='SAMME',
                           base_estimator=DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=1,
                           max_features=None, max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                           splitter='best'),
                           learning_rate=1.0, n_estimators=200, random_state=None)
```

```
In [15]: y_pred = clf.predict(X_test)
```

```
In [16]: y_pred
```

[illegible]

```
In [17]: y_test.head()
```

Out[17]:

	class
19	0
819	1
802	2
528	0
905	0

```
In [18]: print(confusion_matrix(y_test, y_pred))
```

```
[ [254    9    0    0]
  [ 11   48    0    0]
  [  0   16    0    0]
  [  0    8    0    0] ]
```

```
In [19]: print(classification_report(y_test, y_pred))
```


	precision	recall	f1-score	support
0	0.96	0.97	0.96	263
1	0.59	0.81	0.69	59
2	0.00	0.00	0.00	16
3	0.00	0.00	0.00	8
micro avg	0.87	0.87	0.87	346
macro avg	0.39	0.44	0.41	346
weighted avg	0.83	0.87	0.85	346

```
In [20]: accuracy_test_ada=round(clf.score(X_test,y_test)*100,2)
accuracy_train_ada=round(clf.score(X_train,y_train)*100,2)
accuracy_ada=round(accuracy_score(y_test, y_pred)*100,2)
print('Training accuracy of Ada Boosted Decision Tree',accuracy_train_ada)
print('Testing accuracy of Ada Boosted Decision Tree',accuracy_test_ada)
print('Accuracy of Ada Boosted Decision Tree:',accuracy_ada)
```

```
Training accuracy of Ada Boosted Decision Tree 88.32
Testing accuracy of Ada Boosted Decision Tree 87.28
Accuracy of Ada Boosted Decision Tree: 87.28
```

```
In [21]: cm=confusion_matrix(y_test, y_pred)
cm_df = pd.DataFrame(cm,
                      index = ['unacc', 'acc', 'good', 'vgood'],
                      columns = ['unacc', 'acc', 'good', 'vgood'])
```

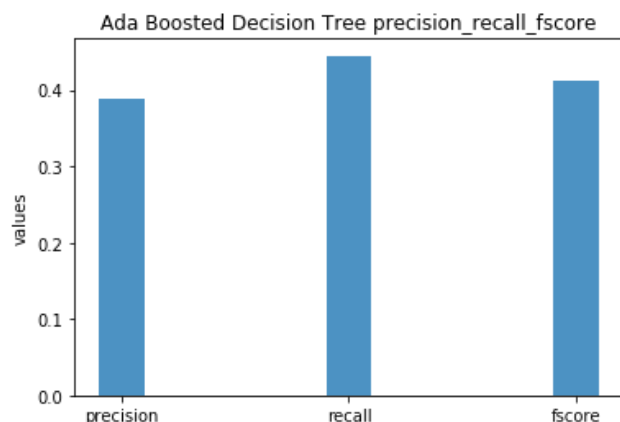
```
In [22]: plt.figure(figsize=(5.5,4))
sns.heatmap(cm_df, annot=True)
plt.title('Ada Boosted Decision Tree Accuracy:{0:.3f}'.format(accuracy_test_ada))
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.show()
```



```
In [23]: pprf = precision_recall_fscore_support(y_test, y_pred, average='macro')
print("Ada Boosted Decision Tree precision_recall_fscore_support ", pprf)
pps = precision_score(y_test, y_pred, labels=None, pos_label=1, average='macro', sample_weight=None)
print("Ada Boosted Decision Tree precision_score -> %.2f"%pps)
prs = recall_score(y_test, y_pred, labels=None, pos_label=1, average='macro', sample_weight=None)
print("Ada Boosted Decision Tree recall_score -> %.2f"%prs)
pfl=f1_score(y_test, y_pred, labels=None, pos_label=1, average='macro', sample_weight=None)
print("f1_score",f1_score(y_test, y_pred, labels=None, pos_label=1, average='macro', sample_weight=None))
print('Ada Boosted Decision Tree f1 score -> %.2f'%pfl)
```

```
Ada Boosted Decision Tree precision_recall_fscore_support (0.3877707896575821,
0.4448346974286267, 0.41195887445887447, None)
Ada Boosted Decision Tree precision_score -> 0.39
Ada Boosted Decision Tree recall_score -> 0.44
f1_score 0.41195887445887447
Ada Boosted Decision Tree f1 score -> 0.41
```

```
In [24]: plt.bar(['precision', 'recall', 'fscore'], [pps, prs, pfl], align='center', alpha=0.8,
width=.2)
plt.ylabel('values')
plt.title('Ada Boosted Decision Tree precision_recall_fscore')
plt.show()
```



```
In [25]: clf = RandomForestClassifier(n_estimators=25)
clf.fit(X_train, y_train)
```

```
Out[25]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
max_depth=None, max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=25, n_jobs=None,
oob_score=False, random_state=None, verbose=0,
warm_start=False)
```

```
In [26]: y_pred = clf.predict(X_test)
```

```
In [27]: y_pred
```

```
Out[27]: array([0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2,
0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 2, 1, 0, 0,
1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 3, 0, 0,
0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 2, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 3, 0, 0, 1, 0, 3, 1, 1, 1,
0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0,
1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1,
0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 3, 0, 0,
0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 2, 0,
0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 0,
0, 0, 0, 0, 0, 1, 3, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0,
0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 3, 0, 0, 1, 0, 0, 0, 0, 0, 0,
1, 1, 0, 0, 0, 2, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0])
```

```
In [28]: y_test.head()
```

```
Out[28]:
```

	class
19	0
819	1
802	2
528	0
905	0

```
In [29]: print(confusion_matrix(y_test, y_pred))
```

```
[[260  3  0  0]
 [ 6 50  2  1]
 [ 0 12  4  0]
 [ 0  3  0  5]]
```

```
In [30]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.98	0.99	0.98	263
1	0.74	0.85	0.79	59
2	0.67	0.25	0.36	16
3	0.83	0.62	0.71	8
micro avg	0.92	0.92	0.92	346
macro avg	0.80	0.68	0.71	346
weighted avg	0.92	0.92	0.91	346


```
In [31]: accuracy_test_rf=round(clf.score(X_test,y_test)*100,2)
accuracy_train_rf=round(clf.score(X_train,y_train)*100,2)
accuracy_rf=round(accuracy_score(y_test, y_pred)*100,2)
print('Training accuracy of RandomForestClassifier:',accuracy_train_rf)
print('Testing accuracy of RandomForestClassifier:',accuracy_test_rf)
print('Accuracy of RandomForestClassifier:',accuracy_rf)
```

```
Training accuracy of RandomForestClassifier: 100.0
Testing accuracy of RandomForestClassifier: 92.2
Accuracy of RandomForestClassifier: 92.2
```

```
In [32]: cm=confusion_matrix(y_test, y_pred)
cm_df = pd.DataFrame(cm,
                      index = ['unacc', 'acc', 'good', 'vgood'],
                      columns = ['unacc', 'acc', 'good', 'vgood'])
```




```
In [33]: plt.figure(figsize=(5.5,4))
sns.heatmap(cm_df, annot=True)
plt.title(' RandomForestClassifier:{0:.3f}'.format(accuracy_test_rf))
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.show()
```



```
In [34]: pprf = precision_recall_fscore_support(y_test, y_pred, average='macro')
print("RandomForestClassifier precision_recall_fscore_support ", pprf)
pps = precision_score(y_test, y_pred, labels=None, pos_label=1, average='macro', sample_weight=None)
print("RandomForestClassifier precision_score -> %.2f"%pps)
prs = recall_score(y_test, y_pred, labels=None, pos_label=1, average='macro', sample_weight=None)
print("RandomForestClassifier recall_score -> %.2f"%prs)
pfl=f1_score(y_test, y_pred, labels=None, pos_label=1, average='macro', sample_weight=None)
print("f1_score",f1_score(y_test, y_pred, labels=None, pos_label=1, average='macro', sample_weight=None))
print('RandomForestClassifier f1 score -> %.2f'%pfl)
```

```
RandomForestClassifier precision_recall_fscore_support (0.8031844316674038, 0.677762695753045, 0.7120776050527625, None)
RandomForestClassifier precision_score -> 0.80
RandomForestClassifier recall_score -> 0.68
f1_score 0.7120776050527625
RandomForestClassifier f1 score -> 0.71
```

```
In [35]: plt.bar(['precision', 'recall', 'fscore'], [pps, prs, pfl], align='center', alpha=0.8,
width=.2)
plt.ylabel('values')
plt.title('RandomForestClassifier precision_recall_fscore')
plt.show()
```



```
In [ ]:
```