

```
In [279]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.metrics import accuracy_score, mean_squared_error, r2_score
from sklearn.model_selection import train_test_split
import statsmodels.api as sm
import seaborn as sns
import statsmodels.formula.api as smf
from sklearn.linear_model import LogisticRegression
```

```
In [280]: #for 3-Dimensional representation
from mpl_toolkits.mplot3d.axes3d import get_test_data
# This import registers the 3D projection, but is otherwise unused.
from mpl_toolkits.mplot3d import Axes3D
```

```
In [281]: #reading data from CSV
data = pd.read_csv("advertising.csv")
```

```
In [282]: #to display the top 5 instances in the dataframe
data.head()
```

Out[282]:

	Unnamed: 0	TV	Radio	Newspaper	Sales
0	1	230.1	37.8	69.2	22.1
1	2	44.5	39.3	45.1	10.4
2	3	17.2	45.9	69.3	9.3
3	4	151.5	41.3	58.5	18.5
4	5	180.8	10.8	58.4	12.9

```
In [283]: #to check the columns present in the dataframe
data.columns
```

Out[283]: Index(['Unnamed: 0', 'TV', 'Radio', 'Newspaper', 'Sales'], dtype='object')

```
In [284]: #dropping the unwanted columns in the dataframe
data=data.drop(['Unnamed: 0'], axis=1)
```

```
In [285]: sns.pairplot(data, x_vars=['TV', 'Radio', 'Newspaper'], y_vars='Sales', size=7, aspect=0.7, kind='reg')
```


/Applications/anaconda2/envs/py36/lib/python3.5/site-packages/seaborn/axisgrid.py:2065: UserWarning: The `size` parameter has been renamed to `height`; please update your code.

warnings.warn(msg, UserWarning)


/Applications/anaconda2/envs/py36/lib/python3.5/site-packages/scipy/stats/stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval

```
Out[285]: <seaborn.axisgrid.PairGrid at 0x1c18510630>
```



```
In [286]: #mentioning the size of the plot/figure.
plt.figure(figsize=(16, 8))
#mentioning the scatter function to consider which columns in the dataframe
plt.scatter(
    data['TV'],
    data['Sales'],
    c='red'
)
#Labeling both X & Y axis's
plt.xlabel("Money spent on TV ads ($)")
plt.ylabel("Sales ($)")
#To display the scatter plot with above mentioned attributes
plt.show()
```



```
In [287]: #reshape allows to change the shape of the dataframe without losing/changing data
X = data['TV'].values.reshape(-1,1)
y = data['Sales'].values.reshape(-1,1)

#getting the Linear Regression function from (from sklearn.linear_model import LinearRegression)
reg = LinearRegression()
#
reg.fit(X, y)
```

```
Out[287]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
                        normalize=False)
```

```
In [288]: #accuracy of sales prediction when only TV add's investments in considered
#Returns the coefficient of determination R^2 of the prediction.
reg.score(X,y)
```

```
Out[288]: 0.611875050850071
```


```
In [289]: print(reg.coef_[0][0])
print(reg.intercept_[0])

print("The linear model is: Y = {:.5} + {:.5}X".format(reg.intercept_[0], reg.coef_[0][0]))

0.047536640433019764
7.032593549127693
The linear model is: Y = 7.0326 + 0.047537X
```

```
In [290]: predictions = reg.predict(X)
plt.figure(figsize=(16, 8))
plt.scatter(
    data['TV'],
    data['Sales'],
    c='red'
)
plt.plot(
    data['TV'],
    predictions,
    c='blue',
    linewidth=2
)

plt.xlabel("Money spent on TV ads ($)")
plt.ylabel("Sales ($)")
plt.show()
```



In [291]: *#different statistical values regarding the model on the dataset.*

```
X = data['TV']
y = data['Sales']

X2 = sm.add_constant(X)
est = sm.OLS(y, X2)
est2 = est.fit()
print(est2.summary()) #prints the entire summary of the dataset.
```

```

                        OLS Regression Results
=====
Dep. Variable:          Sales      R-squared:          0.612
Model:                  OLS       Adj. R-squared:       0.610
Method:                 Least Squares   F-statistic:       312.1
Date:                  Thu, 17 Jan 2019   Prob (F-statistic): 1.47e-42
Time:                  02:00:48    Log-Likelihood:    -519.05
No. Observations:      200          AIC:              1042.
Df Residuals:          198          BIC:              1049.
Df Model:               1
Covariance Type:       nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
const                7.0326      0.458      15.360      0.000      6.130      7.935
TV                   0.0475      0.003      17.668      0.000      0.042      0.053
=====
Omnibus:              0.531    Durbin-Watson:       1.935
Prob(Omnibus):        0.767    Jarque-Bera (JB):     0.669
Skew:                 -0.089    Prob(JB):             0.716
Kurtosis:             2.779    Cond. No.             338.
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [292]: #NOW CONSIDERING INVESTMENTS ON BOTH TV AND NEWSPAPER TO PREDICT SALES, SO REMOVING SALES, UNNAMED A  
ND RADIO  
#FROM THE DATASET, WHICH KEEPS ONLY TV AND NEWSPAPER AS ATTRIBUTES IN 'X'  
  
Xs = data.drop(['Sales', 'Radio'], axis=1)  
Xs.head()  
#y = data['Sales'].reshape(-1,1)  
  
reg1 = LinearRegression()  
reg1.fit(Xs, y)
```

```
Out[292]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,  
normalize=False)
```

```
In [293]: #intercept and co-efficient for multiple regression  
  
print(reg1.coef_)  
print(reg1.intercept_)  
  
[0.04690121 0.04421942]  
5.774947967911631
```

```
In [296]: print("The linear model is: Y = {:.5} + {:.5}*TV + {:.5}*newspaper".format(reg1.intercept_, reg1.coef_  
[0], reg1.coef_[1]))  
  
The linear model is: Y = 5.7749 + 0.046901*TV +0.044219*newspaper
```

```
In [297]: #Returns the coefficient of determination R^2 of the prediction.  
reg1.score(Xs, y)
```

```
Out[297]: 0.6458354938293271
```


```
In [298]: df2 = data.drop(['Radio'], axis=1)
```

```
In [299]: model = smf.ols(formula='Sales ~ TV + Newspaper', data=df2)  
results_formula = model.fit()
```



```
In [300]: x_surf, y_surf = np.meshgrid(np.linspace(data.TV.min(), data.TV.max(), 300), np.linspace(data.Newspaper.min(), data.Newspaper.max(), 100))
onlyX = pd.DataFrame({'TV': x_surf.ravel(), 'Newspaper': y_surf.ravel()})
fittedY = results_formula.predict(exog=onlyX)
```

```
In [301]: fig = plt.figure(figsize=(16, 16))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(data['TV'], data['Newspaper'], data['Sales'], c='red', marker='o', alpha=0.5)
ax.plot_surface(x_surf, y_surf, fittedY.values.reshape(100, 300), color='black', alpha=0.19)
ax.set_xlabel('TV')
ax.set_ylabel('NewsPaper')
ax.set_zlabel('Sales')
plt.show()
```



```
In [302]: data_lgr = pd.read_csv("advertising.csv")
```

```
In [303]: xdfldr = data_lgr.drop(['Unnamed: 0', 'Sales'], axis=1)
xdfldr.head()
```

Out[303]:

	TV	Radio	Newspaper
0	230.1	37.8	69.2
1	44.5	39.3	45.1
2	17.2	45.9	69.3
3	151.5	41.3	58.5
4	180.8	10.8	58.4

```
In [304]: ydflgr=dfldr[ 'Sales' ]  
ydflgr.head()
```

```
Out[304]: 0    22.1  
1    10.4  
2     9.3  
3    18.5  
4    12.9  
Name: Sales, dtype: float64
```

```
In [305]: len(Xlgr)
```

```
Out[305]: 200
```

```
In [306]: X_train,X_test,y_train,y_test=train_test_split(xdflgr,ydflgr,test_size=0.3,random_state=32)
```

```
In [307]: #Logistic Regresion  
log_clf = LogisticRegression(random_state=0, solver='lbfgs',multi_class='multinomial')
```

```
In [ ]:
```